

“If this district were monitored or audited today, what appears satisfied, what is at risk, what evidence supports that conclusion, and what needs to happen next?”

The one question the platform answers — continuously, not once a year.

THE PROBLEM

A district spends federal special-education money under rules it cannot fully see. Maintenance of effort, excess cost, proportionate share, CEIS — each is a separate calculation, drawing on a different system, governed by federal regulation with a state overlay on top. Nobody in the building owns the whole picture.

Districts find out they were non-compliant when a monitor tells them. By then the fiscal year is closed, and the shortfall comes back out of **local** dollars — the Department recovers from the state, the state from the district.

Today's tooling is an inherited spreadsheet plus a once-a-year consultant.

WHAT WE'RE BUILDING

A platform that ingests data from the systems districts already run — SIS, IEP platform, ERP — normalizes it into a canonical model, evaluates versioned regulatory rules against it, and **explains every result down to the source record**.

It is a **system of assurance, not a system of record**. Existing systems stay authoritative. We read from them; we never write back. That constraint is what makes it installable in a district with no appetite for another migration.

WHO IT'S FOR

The **special-education director** who owns the finding, and the **business manager or CFO** who owns the local dollars that come back out if it goes wrong. Later, the other side of the table: **state education agencies** monitoring every LEA they oversee. Selling district-first is what earns the rule corpus that makes the state product credible.

DELIBERATELY NOT BUILDING

Another IEP authoring platform, SIS, ERP, Medicaid billing engine, or generalized workflow builder. No autonomous write-back into district systems. Each trades a narrow, defensible product for a large support burden.

WHY THIS IS HARD — AND WHY THAT'S THE OPPORTUNITY

Special-education fiscal rules don't behave the way accounting intuition expects. Three things a spreadsheet almost never gets right:

- 1 **MOE has two standards, not one.** Eligibility tests next year's *budget*; compliance tests last year's *actual expenditures*. A district can pass one and fail the other in the same year.
- 2 **Four ways to qualify, only one has to pass** — local funds only, state and local combined, and each on a per-child basis. A tool that checks one method reports failures that aren't failures.
- 3 **Failing MOE doesn't reset the baseline.** Under the subsequent-years rule, next year's required level is what *would* have been required absent the failure — not the reduced amount actually spent. A district modeling from what it really spent digs the hole deeper.

Encoding this correctly, per state, with citations and effective dates, is the work. It is also the barrier to entry.

HOW IT'S BUILT

Deterministic rules own every decision. Outcomes come from versioned logic reproducible years later from the same snapshot. AI classifies documents and drafts remediation — advisory only, validated by a person. AI never issues a determination.

Every conclusion has a “Why.” Requirement, citation, rule-pack version, data snapshot, inputs, arithmetic, source file and rows — and which alternative methods pass.

Regulatory logic is content, not code. Federal baseline, state overlay, local policy — layered and versioned, reviewable by an attorney who doesn't read TypeScript. Adding Alabama doesn't fork the product.

History is reproducible. A finalized assessment is immutable. Ask what the platform concluded on a date; it reconstructs that answer exactly.

NOW	NEXT	THEN	LATER
IDEA fiscal — MOE, excess cost, proportionate share, CEIS/CCEIS	Evidence vault, corrective actions, monitoring readiness	Programmatic SPED monitoring — Child Find and evaluation timelines; disproportionality	Title I and federal programs, Part C, Head Start, subrecipient monitoring, Section 504

THE DEFENSIBLE ASSET

Not the dashboard, and not the model. It's the versioned regulatory graph with source provenance, the audited calculator library, the reusable state rule packs, the canonical compliance ontology, the evidence graph, and the integration library. Those compound — a competitor starting today starts at zero on all six. The buyer shifts with them: district assurance tool → multi-state platform → state-agency monitoring.

STATUS

Phase 0. Repository, rule-pack engine, restricted rule DSL and CI baseline are built, with federal IDEA Part B rules encoded and carrying their CFR citations. **Calculators are specified, not yet implemented** — by design, the golden test corpus is written from the statute before the function.